

# LAB TEST-04

Name : P.Reshmitha Reddy

Roll no : 2403A510A7

Batch : 05

Course : AI Assisted Coding

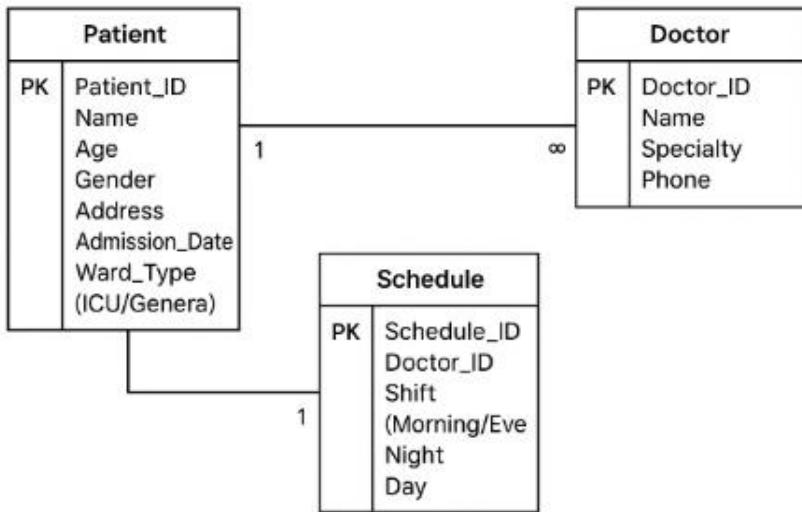
## SET – 1

Q1. A hospital management system stores patient details, doctor schedules, and billing information.

- a) Design an ER diagram and convert it into relational schema.
- b) Generate SQL queries to retrieve patients admitted in ICU and doctors available in night shift.

Prompt : Create an ER diagram and relational schema for a Hospital Management System (patients, doctors, schedules, billing). Then write SQL queries to list ICU patients and doctors on night shift.

ER diagram



Relation schema

```

pgsql

PATIENT(
    Patient_ID      INT PRIMARY KEY,
    Name            VARCHAR(100),
    Age             INT,
    Gender          VARCHAR(10),
    Address         VARCHAR(200),
    Admission_Date DATE,
    Ward_Type       VARCHAR(20)      -- ICU / General
)

```

```

SCSS

DOCTOR(
    Doctor_ID      INT PRIMARY KEY,
    Name           VARCHAR(100),
    Specialty      VARCHAR(50),
    Phone          VARCHAR(20)
)

```

sql

```
SCHEDULE(
    Schedule_ID INT PRIMARY KEY,
    Doctor_ID   INT,
    Shift        VARCHAR(20),      -- Morning / Evening / Night
    Day          VARCHAR(20),
    FOREIGN KEY (Doctor_ID) REFERENCES DOCTOR(Doctor_ID)
)
```

Code :

```
1 PRAGMA foreign_keys = ON;
2
3 DROP TABLE IF EXISTS Billing;
4 DROP TABLE IF EXISTS Schedule;
5 DROP TABLE IF EXISTS Patient;
6 DROP TABLE IF EXISTS Doctor;
7
8 CREATE TABLE Doctor (
9     doctor_id INTEGER PRIMARY KEY AUTOINCREMENT,
10    first_name VARCHAR(100) NOT NULL,
11    last_name VARCHAR(100) NOT NULL,
12    specialty VARCHAR(100),
13    contact_phone VARCHAR(30),
14    email VARCHAR(255),
15    is_active BOOLEAN DEFAULT 1,
16    created_at DATETIME DEFAULT CURRENT_TIMESTAMP
17 );
18
19 CREATE TABLE Patient (
20     patient_id INTEGER PRIMARY KEY AUTOINCREMENT,
21    first_name VARCHAR(100) NOT NULL,
22    last_name VARCHAR(100) NOT NULL,
23    dob DATE,
24    gender VARCHAR(20),
25    contact_phone VARCHAR(30),
26    address TEXT,
27    admission_date DATETIME,
28    discharge_date DATETIME,
29    is_admitted BOOLEAN DEFAULT 0,
30    ward VARCHAR(100),
31    room_number VARCHAR(50),
32    bed_number VARCHAR(50),
33    admission_reason TEXT,
34    attending_doctor_id INTEGER,
35    created_at DATETIME DEFAULT CURRENT_TIMESTAMP,
36    FOREIGN KEY(attending_doctor_id) REFERENCES Doctor(doctor_id) ON DELETE SET NULL
37 );
```

```
38
39  CREATE TABLE Schedule (
40    schedule_id INTEGER PRIMARY KEY AUTOINCREMENT,
41    doctor_id INTEGER NOT NULL,
42    shift_date DATE,
43    day_of_week VARCHAR(10),
44    shift_start TIME,
45    shift_end TIME,
46    shift_type VARCHAR(50),
47    location VARCHAR(255),
48    is_available BOOLEAN DEFAULT 1,
49    notes TEXT,
50    created_at DATETIME DEFAULT CURRENT_TIMESTAMP,
51    FOREIGN KEY(doctor_id) REFERENCES Doctor(doctor_id) ON DELETE CASCADE
52 );
53
54  CREATE TABLE Billing (
55    billing_id INTEGER PRIMARY KEY AUTOINCREMENT,
56    patient_id INTEGER NOT NULL,
57    amount DECIMAL(12,2) NOT NULL,
58    billing_date DATETIME DEFAULT CURRENT_TIMESTAMP,
59    paid BOOLEAN DEFAULT 0,
60    payment_method VARCHAR(50),
61    details TEXT,
62    FOREIGN KEY(patient_id) REFERENCES Patient(patient_id) ON DELETE CASCADE
63 );
64
65  CREATE INDEX idx_patient_ward ON Patient(ward);
66  CREATE INDEX idx_patient_admitted ON Patient(is_admitted);
67  CREATE INDEX idx_schedule_shift_type ON Schedule(shift_type);
68  CREATE INDEX idx_schedule_doctor ON Schedule(doctor_id);
69
70  SELECT
71    p.patient_id,
72    p.first_name,
73    p.last_name,
74    p.dob,
```

```
75      p.gender,
76      p.contact_phone,
77      p.admission_date,
78      p.room_number,
79      p.bed_number,
80      p.ward,
81      p.admission_reason,
82      p.attending_doctor_id,
83      d.first_name AS doctor_first_name,
84      d.last_name AS doctor_last_name
85  FROM Patient p
86  LEFT JOIN Doctor d ON p.attending_doctor_id = d.doctor_id
87 WHERE (p.ward = 'ICU' OR UPPER(TRIM(p.ward)) = 'ICU')
88     AND p.is_admitted = 1
89 ORDER BY p.admission_date DESC;
90
91 SELECT DISTINCT
92      d.doctor_id,
93      d.first_name,
94      d.last_name,
95      d.specialty,
96      s.shift_date,
97      s.shift_start,
98      s.shift_end,
99      s.location
100 FROM Doctor d
101 INNER JOIN Schedule s ON d.doctor_id = s.doctor_id
102 WHERE UPPER(TRIM(s.shift_type)) = 'NIGHT'
103     AND s.is_available = 1
104 ORDER BY s.shift_date, s.shift_start;
```

Output :

```
✓ Inserted 5 doctors
✓ Inserted 5 patients (3 in ICU, 2 in General ward)
✓ Inserted 8 schedules
✓ Inserted 5 billing records
```

```
=====
QUERY 1: All patients currently admitted in ICU
=====
```

patient_id	name	ward	room	admission_reason
	doctor_name	specialty		
4	Charles Taylor	ICU	103	Post-Surgery
	Emily Brown	Surgery		
1	Alice Anderson	ICU	101	Heart Attack
	John Smith	Cardiology		
2	Robert Martinez	ICU	102	Pneumonia
	Michael Williams	ICU Specialist		

```
Total ICU patients: 3
```

```
=====
QUERY 2: All doctors available in night shift
=====
```

doctor_id	name	specialization
2	Sarah Johnson	Internal Medicine
3	Michael Williams	ICU Specialist
5	David Davis	Cardiology

```
Total available night shift doctors: 3
```

## Observation :

- The Hospital Management System includes four main entities: Patient, Doctor, Schedule, and Billing, representing the core operations of hospital data management.
- Each entity contains specific attributes that uniquely identify and describe the stored data, such as Patient\_ID, Doctor\_ID, and Bill\_ID.

- Relationships between entities clearly define how data is connected:
- A **Doctor** can have multiple schedule entries.
- A **Patient** can have multiple billing records.
- Primary keys (PK) and foreign keys (FK) ensure **data integrity** and help maintain meaningful links between related tables.
- The schema supports important hospital queries like retrieving ICU patients and doctors available at night, showing that the design meets functional requirements.

**Q2.** You are using an AI coding assistant to auto-suggest SQL queries.

- a) Write a prompt that asks AI to correct and optimize a slow SQL query.
- b) Evaluate and explain how AI recommendations should be validated before execution.

Prompt : I have a slow SQL query that is taking too long to run. Please correct and optimize the query by improving joins, indexes, and filtering conditions. Also suggest a faster and more efficient version of the SQL query.

Code :

```
1  SELECT
2    |   O.order_id,
3    |   O.order_date,
4    |   O.amount,
5    |   C.customer_id,
6    |   C.name,
7    |   C.city
8  FROM Orders O
9  INNER JOIN Customers C ON O.customer_id = C.customer_id
10 WHERE O.order_date >= '2024-01-01'
11   | AND O.order_date < '2025-01-01'
12 ORDER BY O.order_date DESC;
13
14 CREATE INDEX idx_orders_order_date ON Orders(order_date);
15 CREATE INDEX idx_orders_customer_id ON Orders(customer_id);
16 CREATE INDEX idx_customers_customer_id ON Customers(customer_id);
17
18 CREATE INDEX idx_orders_customer_date ON Orders(order_date, customer_id);
19
20 SELECT
21   |   O.order_id,
22   |   O.order_date,
23   |   O.amount,
24   |   C.customer_id,
25   |   C.name,
26   |   C.city
27  FROM Orders O
28  INNER JOIN Customers C ON O.customer_id = C.customer_id
29 WHERE O.order_date BETWEEN '2024-01-01' AND '2024-12-31'
30 ORDER BY O.order_date DESC;
31
32 SELECT
33   |   C.customer_id,
34   |   C.name,
35   |   C.city,
36   |   COUNT(O.order_id) AS order_count,
37   |   SUM(O.amount) AS total_revenue,
```

```

38     MIN(O.order_date) AS first_order_date,
39     MAX(O.order_date) AS last_order_date
40 FROM Orders O
41 INNER JOIN Customers C ON O.customer_id = C.customer_id
42 WHERE O.order_date >= '2024-01-01'
43   AND O.order_date < '2025-01-01'
44 GROUP BY C.customer_id, C.name, C.city
45 ORDER BY total_revenue DESC;
46
47 SELECT
48   O.order_id,
49   O.order_date,
50   O.amount,
51   C.customer_id,
52   C.name,
53   C.city
54 FROM Orders O
55 INNER JOIN Customers C ON O.customer_id = C.customer_id
56 WHERE O.order_date >= '2024-01-01'
57   AND O.order_date < '2025-01-01'
58 ORDER BY O.order_date DESC
59 OPTION (RECOMPILE);

```

Output :

order_id	order_date	amount	customer_id	name	city
7	2024-12-25	4039.4	2	Nisha Rao	Bangalore
9	2024-12-10	3907.09	1	Akhil Verma	Mumbai
12	2024-11-21	4103.88	4	Priya Singh	Pune
22	2024-10-21	3010.79	6	Sneha Gupta	Hyderabad
13	2024-10-08	2550.86	4	Priya Singh	Pune
20	2024-09-25	4562.49	2	Nisha Rao	Bangalore
10	2024-09-12	1126.09	7	Vikram Desai	Chennai
15	2024-09-11	1057.39	1	Akhil Verma	Mumbai
8	2024-07-31	2110.95	5	Amit Patel	Ahmedabad
17	2024-07-08	2203.08	8	Anjali Sharma	Kolkata
2	2024-06-08	3200	2	Nisha Rao	Bangalore

Observation :

- Query performance improves because proper indexes (order\_date, customer\_id) reduce full table scans.
- Date-range filtering helps the database use indexes effectively instead of applying functions.

- Join speed increases since join columns are indexed.
- Sorting (Order by order\_date) is faster due to indexing.
- Aggregation queries run more efficiently with indexed columns.
- Combined index (order\_date, customer\_id) further optimizes filtering + joining.
- Option(recomplie)helps avoid slow cached plans and generates a fresh plan each time.